REFINERY

bank_account_original.cob

Independent Verification & Validation

AI-Modified: attepmt_2_germini.cob

Reference: REF-20260513-6268

Date: 2026-05-13

Client: DEMO-BANK-2026

Refinery 0.1.0

PASS

1. EXECUTIVE SUMMARY

Semantically Equivalent - Approved for Deployment

No semantic divergences detected. The AI-modified program is functionally identical to the original.

CPU time reduction	-3.1%
Semantic checks run	6
Divergences found	0
Risk score	**0/100**
Output equivalence	**SKIPPED**

2. METHODOLOGY

CPU Estimation: Both programs run through Refinery's SyntheticRunner, which models instruction-level cost on z/Architecture. Estimates are deterministic and calibrated to hardware measurements.

Semantic Analysis — 6 independent checks: (1) AST-based COMPUTE expression comparison (HIGH if any expression changes). (2) PIC data type comparison (HIGH for size changes; HIGH for COMP-3 additions — packed decimal changes can cause 0C7 abends if downstream systems expect display format). (3) PERFORM/SEARCH/SORT verb counts (MEDIUM if counts differ). (4) REDEFINES clause offset analysis — detects offset or length shifts in WORKING-STORAGE overlay fields, a common source of silent data corruption after AI edits (HIGH). (5) Procedure call graph diff — detects reordered, dropped, or added PERFORM edges (HIGH/MEDIUM). (6) GnuCOBOL output equivalence — when available, both programs are compiled and executed with synthetic test inputs; stdout divergence is flagged OUTPUT_DIVERGENCE HIGH.

Verdict Criteria: PASS requires all checks clean. Any HIGH or MEDIUM finding blocks deployment. LOW-only findings permit conditional approval.

3. EXECUTION EVIDENCE

CPU Performance

Metric	Original	AI-Modified	Change
CPU time (microseconds)	1,181,212	1,218,390	+3.1%

AST Feature Comparison

Feature	Original	AI-Modified
Lines of code	331	337
COMPUTE verbs	1	1

PERFORM nesting depth	15	15
SEARCH verbs	16	16
SORT verbs	0	0
Working storage (bytes)	117	117
File section entries	1	1

Maintainability

Metric	Original	AI-Modified	Change
Cyclomatic complexity (McCabe)	38	38	0

Output Equivalence

SKIPPED — GnuCOBOL not available — equivalence check skipped.

4. DIVERGENCE LOG

No divergences detected. All semantic checks passed.

5. VERDICT AND SIGN-OFF

PASS

The AI-modified program **attepmt_2_germini.cob** is semantically equivalent to **bank_account_original.cob** under the test conditions described in Section 2. No semantic divergences were detected. Deployment decision remains with the authorising engineer.

Reference: REF-20260513-6268 | Date: 2026-05-13 | Refinery: 0.1.0 | Client: DEMO-BANK-2026

This report was generated automatically by Refinery 0.1.0. CPU performance estimates are produced via SyntheticRunner and are indicative only. Semantic correctness analysis (checks 1–6) is performed via static AST analysis and is independent of the runtime environment. Runtime behaviour under z/OS Language Environment, VSAM, or DB2 is outside the scope of this report. This report certifies semantic equivalence only — that the AI-modified program produces identical observable behaviour to the original under the test conditions described in Section 2. It does not certify correctness of business logic, compliance with regulatory requirements, or runtime behaviour under any production environment. Liability is limited to the terms of the applicable engagement agreement. This report does not constitute legal, regulatory, or financial advice.